

CSC 665 Spring 2020 HW2*

Problem 1. This problem is about anytime the Hoeffding's inequality.

(a) As a warmup, derive both versions of the inequalities. That is, let $\hat{\mu}_t = \frac{1}{t} \sum_{s=1}^t X_s$ where X_s is an i.i.d. $(\sigma^2 = 1)$ -sub-Gaussian. Prove the following using first principles (i.e., do not use Hoeffding's bound + variable change):

$$\text{(fixed time)} \quad \text{Fix } t \geq 1. \quad \mathbb{P} \left(\hat{\mu}_t \geq \sqrt{\frac{2}{t} \ln(1/\delta)} \right) \leq \delta$$

$$\text{(anytime)} \quad \mathbb{P} \left(\forall t \geq 1, \hat{\mu}_t \geq \sqrt{\frac{2}{t} \ln(4t^2/\delta)} \right) \leq \delta$$

(b) Your classmate claims that in fact the fixed time deviation bound above actually works for all time t throughout, with probability at least $1 - \delta$. Let us empirically show that she is wrong! Let $N = 10,000$ and $\delta = 0.1$. Draw N Gaussian random variables from mean 0 and variance 1; call them X_1, \dots, X_N . With those, compute $\{\hat{\mu}_t\}_{t=1}^N$. Record whether there exists $t \in [1, N]$ such that $\hat{\mu}_t$ that cross the deviation $\sqrt{(2/t) \ln(1/\delta)}$. Let $Y = 1$ if this was true and 0 otherwise. Now, repeat this 100 times with a fresh set of random samples. Denote by Y_1, \dots, Y_{100} those binary values. If your friend is correct, we must be seeing that the average of $\{Y_i\}_{i=1}^{100}$ is around δ or less.

- Use your favorite programming language to perform the simulation.
- Report the average of $\{Y_i\}$ for both the fixed time version and anytime version.
- Pick some of the random trial and plot $t \cdot \hat{\mu}_t$ and the both deviation bounds multiplied by t , all three of them in one plot with x-axis being t . (Multiplying t is merely to improve the visual).
- Submit your code, plot, and explanations.

*v1

Problem 2. This is about UCB and LinUCB. In the class, we learned a fixed budget version of UCB where we need to feed n , the time horizon, to the algorithm. In this homework, let us use the anytime version of UCB, which selects arms by

$$A_t = \arg \max_{i \in \{1, \dots, k\}} \hat{\mu}_i(t-1) + \sqrt{\frac{2 \log(t^{2.1})}{T_i(t-1)}}$$

where we enforce the objective function to be ∞ when the count $T_i(t-1)$ is zero.

Recall that we assume that the arms $a \in \mathcal{A}$ satisfies that $\|a\|_2 \leq 1$ and $\|\theta^*\|_2 \leq 1$. Let us now relax the latter assumption to $\|\theta^*\|_2 \leq S$ where $S > 0$ is known to us. My apologies for not being exact with the definition of $\sqrt{\beta_t}$ in the class, which contained $\|\theta^*\|_2$. This information is not known to the learner. Instead, we need to use:

$$\sqrt{\beta_t} = \sqrt{\lambda} \cdot S + \sqrt{\log \left(\frac{|V_{t-1}|}{|V_0| \cdot \delta^2} \right)}$$

For other details, please look at the lecture whiteboard shared in piazza (dropbox link).

(a) Implement UCB.

(b) Before implementing LinUCB, derive sequential updates so we have the per-time-step time complexity of $O(d^2)$ w.r.t. the dimension d . Specifically,

- Derive an update equation for V_t^{-1} based on V_{t-1}^{-1} directly (rather than computing the inversion). Use Woodbury matrix identity (see Wikipedia) and the fact that $V_t = V_{t-1} + A_t A_t^\top$.
- Derive an update equation for $|V_t|$ from $|V_{t-1}|$, without directly computing the determinant. Hints can be found somewhere in the lecture whiteboard shared in piazza (dropbox link).

(c) Implement LinUCB. Again, ensure that the per-time-step time complexity must be $O(d^2)$ w.r.t. d . Otherwise, points will be deducted.

(d) Design simulation setups where there are k arms, each with known feature vectors $a \in \mathbb{R}^d$. Compare UCB and LinUCB w.r.t. the cumulative regret. Suggestions:

- Design at least two settings: one where LinUCB might perform better and one where UCB might perform better.
- Submit your code, plots, and explanations.

Don't worry if your designed setting does not work as you expected; this is an open question. Just be sure to provide your thoughts in the answer.

Tip: I emphasized this in the last homework too, but please do spend some time to develop test cases, visually inspect your code, printout values, use step-by-step debugger to check values, etc., to convince yourself that the algorithm is correct. Mathematical code is hard to debug, but the cost of bugs is tremendous.